LOAD TEST REPORT

DATE: 6/24/2014

TEST FROM : VIRGINIA

Query URL: http://69.55.49.79/

Started at: Tue Jun 24 2014, 03:34:35 -04:00

Finished at: Tue Jun 24 2014, 03:34:35 -04:00

ANALYSIS

This rush generated **881** successful hits in **60 seconds** and we transferred **6.67 MB** of data in and out of your app. The average hit rate of **15/second** translates to about **1,268,640** hits/day.

The average response time was **395 ms**.

RESPONSE TIMES

FASTEST: 65 MS

SLOWEST: 997 MS

AVERAGE: 395 MS

TEST CONFIGURATION

REGION: VIRGINIA

DURATION: 60 SECONDS

LOAD: 1-50 USERS

OTHER STATS

AVG. HITS: 15 /SEC

DATA TRANSFERED: 6.67 MB

HITS 89.62% (881)

ERRORS 0.00% (0)

TIMEOUTS 10.38% (102)

HITS

This rush generated **881** successful hits. The number of hits includes all the responses listed below. For example, if you only want **HTTP 200 OK** responses to count as Hits, then you can specify **--status 200** in your rush.

CODE	TYPE	DESCRIPTION	AMOUNT
200	HTTP	OK	881

HTTP 200 OK **100%** (881)

TIMEOUTS

The first timeout happened at **55 seconds** into the test when the number of concurrent users was at **46**. Looks like you've been rushing with a timeout of **1000 ms**. Timeouts tend to increase with concurrency if you have lock contention of sorts. You might want to think about in-memory caching using redis, memcached or varnish to return stale data for a period of time and asynchronously refresh this data.

RESPONSE TIMES

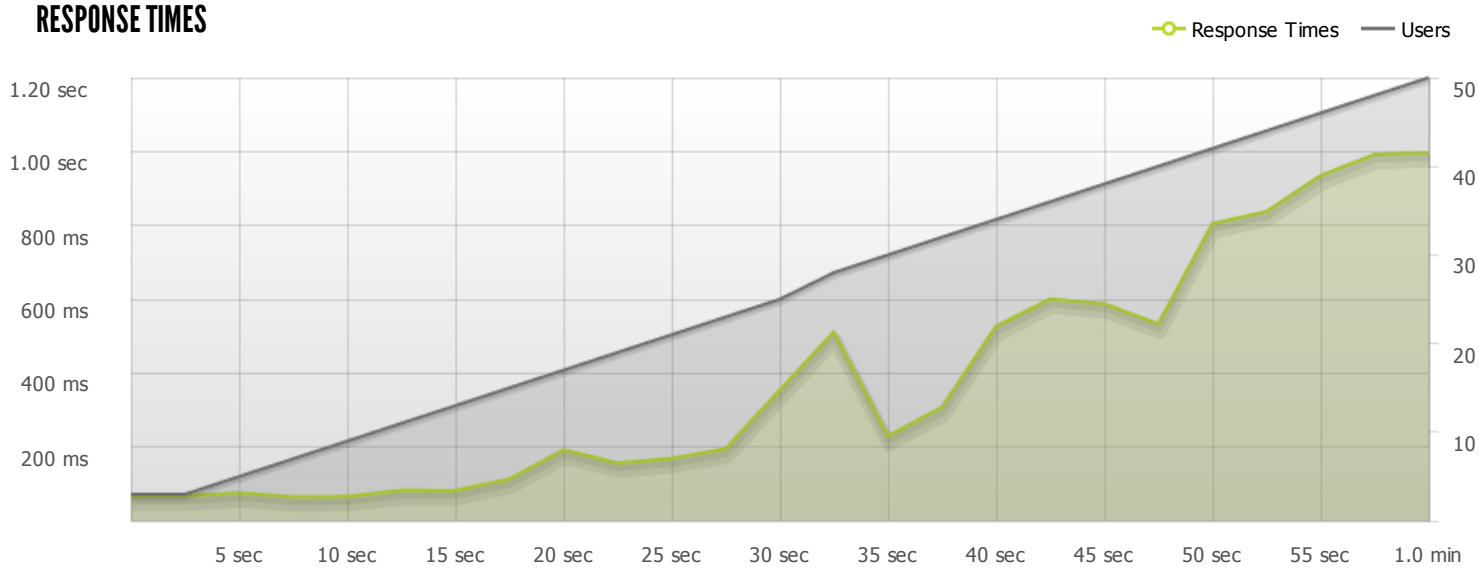

The max response time was: **996 ms @ 50 users**

HIT RATE

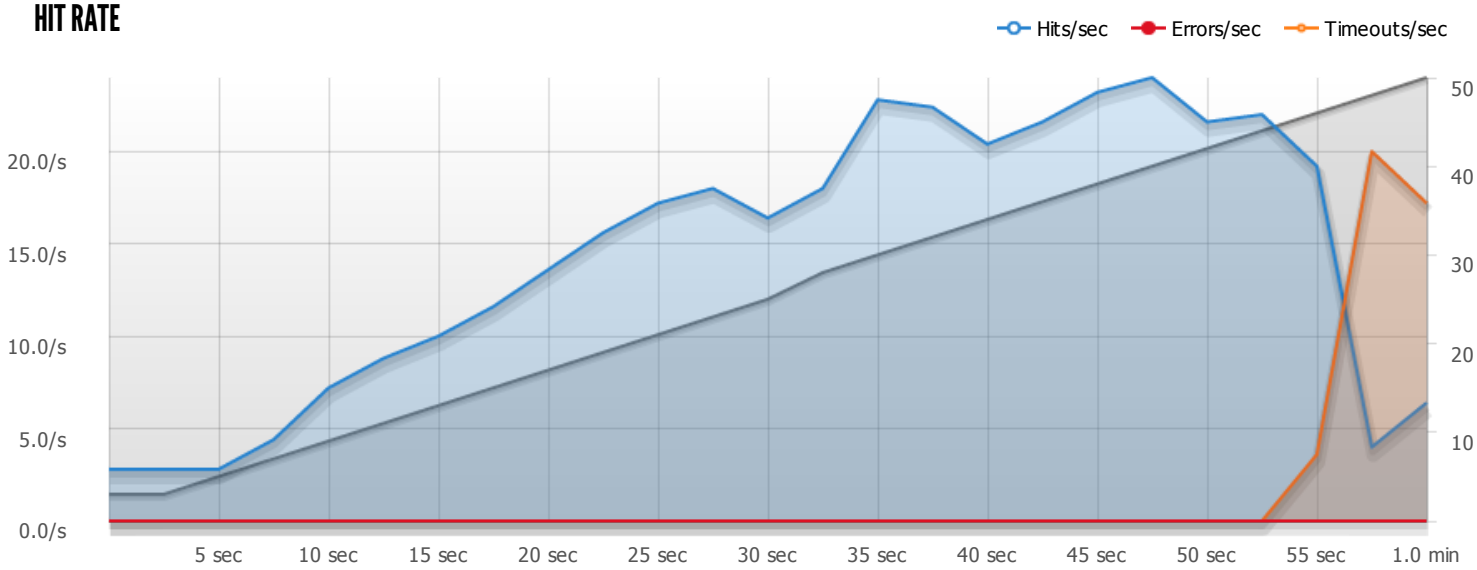

The max hit rate was: **24 hits per second**
