

LOAD TEST REPORT

DATE: 6/26/2014

TEST FROM : VIRGINIA

Query URL: http://69.55.49.190/

Started at: Thu Jun 26 2014, 05:25:44 -04:00

Finished at:

ANALYSIS

This rush generated **680** successful hits in **60 seconds** and we transferred **5.83 MB** of data in and out of your app. The average hit rate of **11/second** translates to about **979,200** hits/day.

The average response time was **404 ms**.

RESPONSE TIMES	TEST CONFIGURATION	OTHER STATS
FASTEST: 96 ms	REGION: VIRGINIA	AVG. HITS: 11 /SEC
SLOWEST: 936 ms	DURATION: 60 SECONDS	DATA TRANSFERED: 5.83 MB
AVERAGE: 404 ms	LOAD: 1-50 USERS	



HITS **73.75%** (680)
ERRORS **0.00%** (0)
TIMEOUTS **26.25%** (242)

HITS

This rush generated **680** successful hits. The number of hits includes all the responses listed below. For example, if you only want **HTTP 200 OK** responses to count as Hits, then you can specify **--status 200** in your rush.

CODE	TYPE	DESCRIPTION	AMOUNT
200	HTTP	OK	680




HTTP 200 OK **100%** (680)

TIMEOUTS

The first timeout happened at **32.5 seconds** into the test when the number of concurrent users was at **28**. Looks like you've been rushing with a timeout of **1000 ms**. Timeouts tend to increase with concurrency if you have lock contention of sorts. You might want to think about in-memory caching using [redis](#), [memcached](#) or [varnish](#) to return stale data for a period of time and asynchronously refresh this data.


RESPONSE TIMES



STEP 1
Response Times

The max response time was: **936 ms @ 44 users**

HIT RATE



STEP 1

Hits/sec Errors/sec Timeouts/sec

The max hit rate was: **23 hits per second**